

Digital-Twin–Driven Semi-Supervised GAN Federated Learning

Sugeet Sood
Vibhor Barguje

18 May 2025

Outline

- 1 Introduction
- 2 Background
- 3 Baseline Sampling
- 4 System Architecture
- 5 Model Details
- 6 Loss Functions
- 7 Client Algorithms
- 8 Client Selection
- 9 Experimental Setup
- 10 Implementation Details
- 11 Future Work

Motivation

- Edge devices generate sensitive data requiring local processing.
- Centralized training compromises privacy and incurs latency.
- Federated Learning (FL) offers a solution but faces non-IID data and device heterogeneity.
- Semi-Supervised GANs (SGANs) leverage unlabeled data but need stabilization.

- 1 **FedProx Integration:** Proximal regularisation for heterogeneity handling.
- 2 **SGAN Enhancements:** Feature Dropout and Mini-Batch Discrimination for robust training.
- 3 **Loss-Guided Selection:** Prioritize clients by loss with 10% exploration.

Federated Learning Essentials

- Clients compute local updates Δw_k on private data.
- Server aggregates: $w^{(t+1)} = w^{(t)} + \sum_k \frac{n_k}{\sum_j n_j} \Delta w_k$ [4].
- Challenges: statistical heterogeneity, system heterogeneity, privacy.

Semi-Supervised GAN (SGAN)

- Discriminator outputs K real classes + 1 fake class [?].
- Combines supervised loss on scarce labels with unsupervised real/fake losses.
- Feature Matching aligns real and generated features.

- **FedProx:** Adds $\frac{\mu}{2} \|\theta - \theta^{(t)}\|^2$ to local objective [5].
- **Loss-Guided Selection:** Rank by last-round loss arl_D , reserve 10% for random exploration [?].

Random Client Sampling Algorithm

basicstyle= for t in $1..T$:

$S_t \leftarrow \text{randomsubsetofclientsbroadcast } w^{(t)}$ parallel for k in S_t : $w_k =$
 $\text{LocalTrain}(w^{(t)})$ $\text{delta}_{w_k} = w_k - w^{(t)}$ $w^{(t+1)} = w^{(t)} + \text{Average}(\text{delta}_{w_k})$

Simple but ignores data utility and client speed.

Layered Architecture

Edge Server

- Selection Module: random / loss-guided.
- Aggregation Engine: FedAvg / FedProx.
- Monitoring: Prometheus Grafana.

Clients

- Local SGAN + FedProx training.
- Metrics reporting: losses, samples.
- Fault handling: timeouts, NaNs.

Digital Twin: Flask webapp for live inference updates.

Generator Architecture

- Input: $z \sim \mathcal{N}(0, I)^{100}$.
- MLP: 100-512-ReLU, 512-1024-BN-ReLU, 1024-3072-Tanh.
- Output: $3 \times 32 \times 32$ image in $[-1,1]$.

Discriminator Architecture

- FC 3072-512-256 with LeakyReLU.
- Feature Dropout ($p=0.5$).
- Mini-Batch Discrimination: 50 kernels.
- Output: $(256 + 50) \rightarrow K + 1$ logits.

Loss Components

- Supervised: smoothed CE on labels.
- Unsupervised Real: $-\log(1 - p_{fake}(x))$.
- Unsupervised Fake: $-\log p_{fake}(G(z))$.
- Adversarial (G): $-\log(1 - p_{fake}(G(z)))$.
- Feature Matching: $\|E[\phi(x)] - E[\phi(G(z))]\|^2$.
- FedProx: $\frac{\mu}{2}\|\theta - \theta^{(t)}\|^2$.

Local Training Pseudocode

```
basicstyle= function LocalTraining(wG,wD): G.load(wG); D.load(wD)  
G0,D0 = G.clone(), D.clone() for epoch=1..E: for batch in data: D  
update: real, sup, fake + prox D.backward(); D.step() G update: adv +  
FM + prox G.backward(); G.step() return G.weights()-wG,  
D.weights()-wD Robust to NaNs and timeouts.
```

Loss-Guided Selector

basicstyle= function SelectClients(hist,M,alpha): $\text{explo} = \text{ceil}(\alpha * M)$
 $\text{explorees} = \text{random}_c \text{choose}(\text{explo})$ $\text{ranked} = \text{sort}(\text{hist.loss}, \text{desc})$ $\text{selected} =$
 $\text{ranked}[: M - \text{explo}]$ $\text{return} \text{explorees} \cup \text{selected}$ Balances exploitation of
high-loss clients with exploration.

Dataset Splits

- CIFAR-10 subset: 2 000 images (seed=42).
- Client 0: 100 labeled; Clients 1–3: 633 unlabeled each.
- Normalized to $[-1,1]$; no augmentation.

Key Metrics

- **Convergence:** 1.3x faster reduction in discriminator loss vs. random sampling.
- **Generation Quality:** 7.4-point FID improvement over random sampling.
- **Efficiency:** Comparable communication per round.

- **DataLoader & Splits:** CIFAR-10 → 2 000 samples, fixed seed, 100 labeled / 633 unlabeled.
- **Client Class:** `get_data_loader()`, `has_labels()`, empty-dataset fallbacks.
- **Reproducibility:** torch/NumPy seed, GPU/CPU toggle, DataLoader workers.

Training Metrics Tracking

- **Hyperparameters:** smoothing=0.1, clip_D=1.0, clip_G=1.0, updates_per_D=1.
- **Client-Returned Metrics:** avg_loss_D, avg_loss_G, samples_processed, epochs_performed.
- **Server Graphing:** accumulate server-side avg losses per round; plot at end.
- **Sample Generation:** fixed noise \rightarrow 8 images each round for visual monitoring.

Future Work

- ① Scale-up: 50+ distributed clients in WAN settings.
- ② Privacy: DP-SGD and secure aggregation integration.
- ③ Optimization: Efficient MBD on IoT hardware.
- ④ Deployment: Real-world sensor network applications.

References I

- [1] I. Goodfellow *et al.*, “Generative Adversarial Nets,” NeurIPS, 2014.
- [2] T. Salimans *et al.*, “Improved Techniques for Training GANs,” NeurIPS, 2016.
- [3] N. Srivastava *et al.*, “Dropout: A Simple Way to Prevent Neural Networks from Overfitting,” JMLR, 2014.
- [4] B. McMahan *et al.*, “Communication-Efficient Learning of Deep Networks from Decentralized Data,” AISTATS, 2017.
- [5] T. Li *et al.*, “Federated Optimization in Heterogeneous Networks,” MLSys, 2020.